

# A Distributed Timing Analysis Framework for Large Designs

Tsung-Wei Huang and Martin D. F. Wong (ECE Dept., UIUC, IL)  
Debjit Sinha, Kerim Kalafala, and Natesan Vankateswaran (IBM, NY)



ECE ILLINOIS

 ILLINOIS


# Distributed Timing – Motivation and Goal

## ❑ Motivation


- ❑ Ever increasing design complexity
  - Hierarchical partition
  - Abstraction
  - Multi-threading timing analysis
- ❑ Too costly to afford high-end machines

## ❑ Create a distributed timing engine

- ❑ Explore a feasible framework
- ❑ Prototype a distributed timer
- ❑ Scalability
- ❑ Performance



Multi-threading in a single machine



Distributed computing on a machine cluster

# State-of-the-art Distributed System Packages

---

- ❑ Open-source cloud computing platforms (<https://hadoop.apache.org/>)
  - ❑ Hadoop
    - Reliable, scalable, distributed MapReduce platform on HDFS
  - ❑ Cassandra
    - A scalable multi-master database with no single points of failure
  - ❑ Chukwa
    - A data collection system for managing large distributed systems
  - ❑ Hbase
    - A scalable, distributed database that supports structured data storage
  - ❑ Zookeeper
    - Coordination service for distributed application
  - ❑ Mesos
    - A high-performance cluster manager with scalable fault tolerance
  - ❑ Spark
    - A fast and general computing engine for iterative MapReduce

# The Questions Are

---

- ❑ ***Are these packages really suitable for our applications?***
  - ❑ Google/Hadoop MapReduce programming paradigm
  - ❑ Spark in-memory iterative batch processing
  
- ❑ ***What are the potential hurdles for EDA to use big-data tools?***
  - ❑ Big-data tools are majorly written in JVM languages
  - ❑ EDA applications highly rely on high-performance C/C++
  - ❑ Rewrites of numbers of codes
  
- ❑ ***What are the differences between EDA and big data?***
  - ❑ Computation intensive vs Data intensive
  - ❑ EDA data is more connected than many of social network

# An Empirical Experiment on Arrival Time Propagation

## ❑ Benchmark


❑ Timing graph from ICCAD 2016 CAD contest (*superblue18*)

- 2.5M nodes
- 3.5M edges

## ❑ Implementation

- ❑ Spark – 4 cores
- ❑ Java, Scala, etc. – 1 core
- ❑ C++ – 1 core

Runtime comparison on arrival time propagation



Implementation	Spark 1.4 (RDD + GraphX Pregel)	Scala (Dijkstra)	C++ (Dijkstra)
Runtime (s)	68.45	10.68	1.50

Overhead of GraphX and message passing

Overhead of JVM



# Key Components in our Framework

---

- ❑ **Multiple-program multiple-data paradigm**

- ❑ Different programs for clients and server
- ❑ Better scalability and work distribution

- ❑ **Non-blocking socket IO**

- ❑ Program returns to users immediately
- ❑ Overlap communication and computation

- ❑ **Event-driven environment**

- ❑ Callback for message read/write events
- ❑ Persistent in memory for efficient data processing

- ❑ **Efficient messaging interface**

- ❑ Network see bytes only
- ❑ Serialization and deserialization of timing data

# Non-blocking IO and Event-driven Loop with Libevent

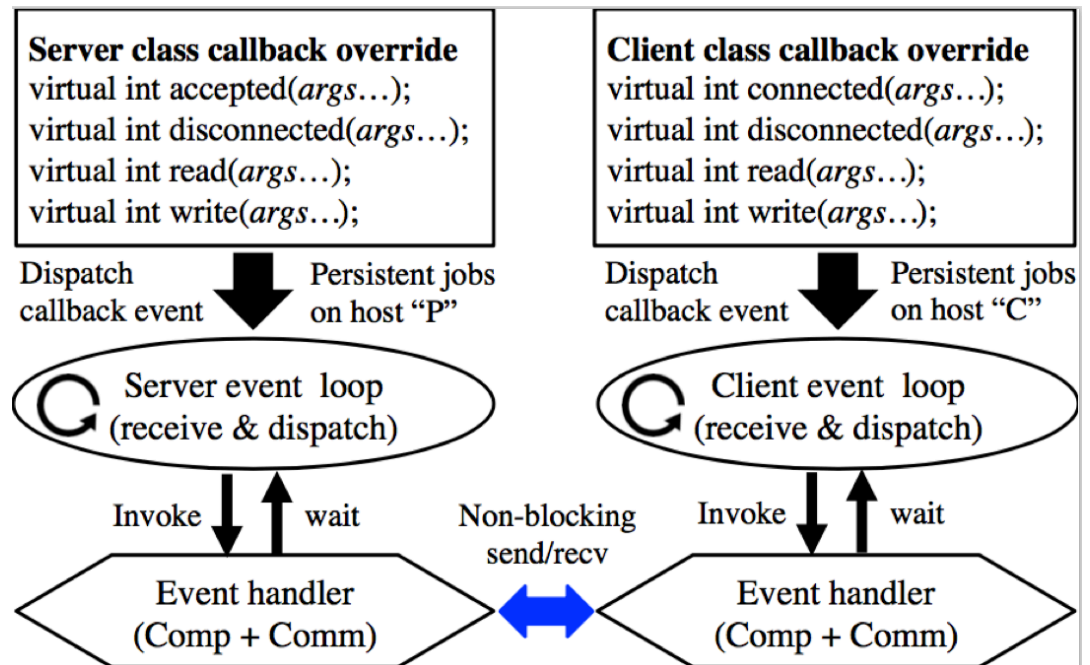
## ❑ Libevent (<http://libevent.org/>)

- ❑ Open-source under BSD license
- ❑ Actively maintained
- ❑ C-based library
- ❑ Non-blocking socket
- ❑ Reactor model



```
// Magic inside dispatch call
while (!event_base_empty(base)) {
    // non-block IO by OS kernel
    active_list ← get_active_events
    foreach(event e in active_list) {
        invoke the callback for event e
    }
}
```

*An example event-driven code*



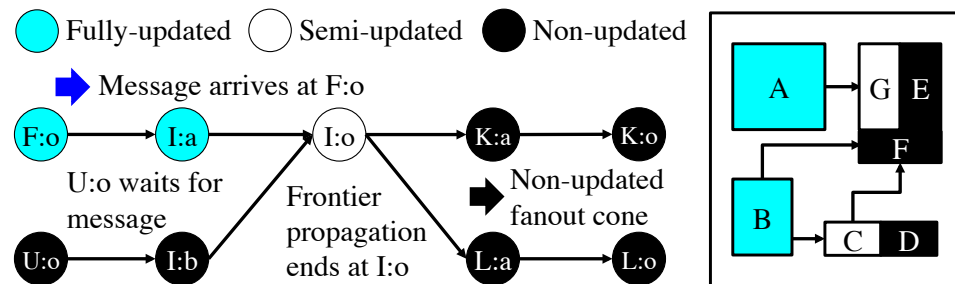
*Interface class in our framework (override virtual methods for event callback)*



# Callback Implementation

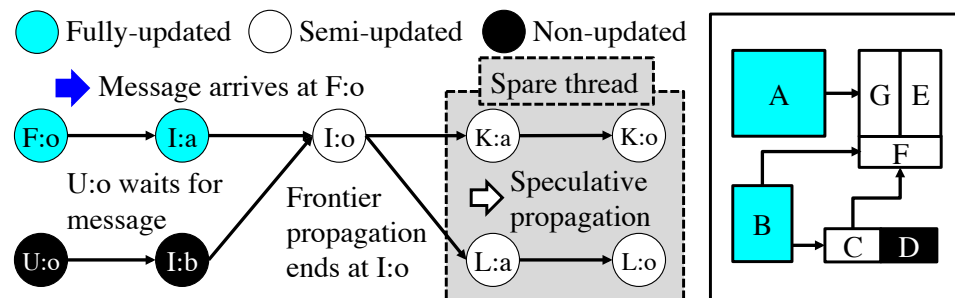
## ❑ Client read callback

- ❑ Receive boundary timing
- ❑ Propagate timing
- ❑ Send back to the server



## ❑ Server read callback

- ❑ Keep boundary mapping
- ❑ Receive boundary timing
- ❑ Propagate timing
- ❑ Send to the client



## ❑ Timing propagation

- ❑ Frontier vs Speculative

*If multi-threading is available, spare thread performs speculative propagation in order gain advanced saving of frontier work*

# Efficient Messaging Interface based on Protocol Buffer

## ❑ Message passing

- ❑ Expensive
- ❑ TCP byte stream
- ❑ Unstructured

## ❑ Data conversion

- ❑ Serialization
- ❑ Deserialization

## ❑ Protocol buffer

- ❑ Customized protocol
- ❑ Simple and efficient
- ❑ Built-in compression

Structured message format  
(.proto)

```
enum KeyType {PIN_NAME}  
enum ValueType {AT, SLACK}  
message Key {  
  optional KeyType type = 1;  
  optional string data = 2;  
}  
message Value {  
  optional ValueType type = 1;  
  optional string data = 2;  
}
```

Google Protocol Buffer  
(open-source compiler)

C++/Java/Python  
source code generator

**.cpp/.h class methods**  
ParseFromArray(void\*, size\_t)  
SerializeToArray(void\*, size\_t)

Message wrapper

**Derived packet struct**  
header\_t header  
void\* buffer

*Integration of Google's open-source protocol buffer into our messaging interface greatly facilitates the data conversion between application-level developments and socket-level TCP byte streams.*

# Evaluation – Software and Hardware Configuration

---

- ❑ **Written in C++ language on a 64-bit Linux machine**
- ❑ **3<sup>rd</sup>-party library**
  - ❑ Libevent for event-driven network programming
  - ❑ Google's protocol buffer for efficient messaging
- ❑ **Benchmarks**
  - ❑ 250 design partitions generated by IBM EinsTimer
  - ❑ Millions-scale graphs generated by TAU and ICCAD contests
- ❑ **Evaluation environment**
  - ❑ UIUC campus cluster (<https://campuscluster.illinois.edu/>)
  - ❑ Each machine node has 16 Intel 2.6GHz cores and 64GB memory
  - ❑ 384-port Mellanox MSX6518-NR FDR InfiniBand (gigabit Ethernet)
  - ❑ Up to 250 machine nodes

# Evaluation – Results and Performance

## ❑ Overall performance

Circuit	$ G $	$ N $	$ V $	$ E $	$ P $	$L$	W/o speculation				W/ speculation			
							cpu	mem	msg	usage	cpu	mem	msg	usage
DesignA	2.2M	1.1M	7.3M	12.4M	250	436	63s	1.6GB	0.7MB	17.3%	76s	1.7GB	1.6MB	64.2%
DesignB	14.5M	9.3M	39.0M	117.0M	37	3216	392s	2.9GB	2.0MB	9.1%	346s	3.1GB	5.7MB	73.1%
DesignC	23.3M	11.3M	76.9M	107.0M	30	2023	478s	4.7GB	2.3MB	19.5%	473s	4.8GB	8.1MB	57.8%
DesignD	42.7M	20.8M	128.1M	178.4M	50	5741	1239s	5.1GB	4.9MB	20.1%	1107s	5.1GB	9.7MB	69.4%

$|G|$ : # of gates.  $|N|$ : # of nets.  $|V|$ : # of nodes.  $|E|$ : # of edges.  $|P|$ : # of partitions.  $L$ : # of levels. cpu: runtime. mem: peak memory on a program. msg: amount of message passing. usage: avg cpu utilization on a program.

## ❑ Scalability

- ❑ Scale to 250 machines (DesignA)

## ❑ Runtime efficiency

- ❑ Less than 1 hour on large designs (DesignC and DesignD)

## ❑ Memory usage

- ❑ Peak usage is only about 5GB on a machine (DesignD)

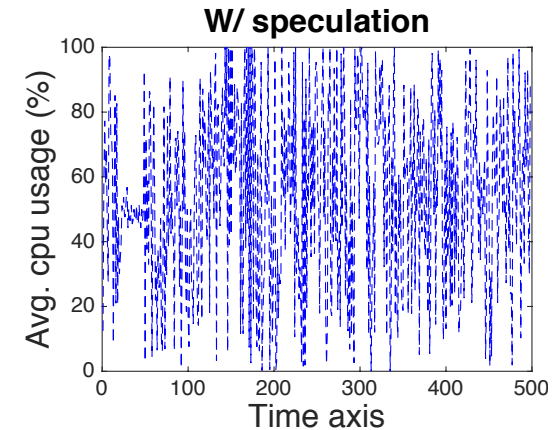

# Evaluation – A Deeper Look

## ❑ CPU utilization

- ❑ W/o speculation
- ❑ W speculation

*W/ speculation on DesignD*  
*+49% cpu rate*


*+4.8MB on message passing*



*Average cpu utilization over time across all machines.*

## ❑ Runtime profile

- ❑ 7% event polling
- ❑ 3% streaming
- ❑ 23% initialization
- ❑ 54% propagation
- ❑ 12% communication



*Runtime profile of our framework (12% on communication and 88% on computation)*

# Conclusion and Future Work

---

## ☐ **Prototype a distributed timing analysis framework**

- ☐ Server-client model
- ☐ Non-blocking socket IO (overlap communication and computation)
- ☐ Event-driven loop (autonomous programming)
- ☐ Efficient messaging interface (serialization and deserialization)

## ☐ **Future work**

- ☐ A system for distributed timing analysis
- ☐ Fault tolerance
- ☐ Distributed common path pessimism removal (CPPR)

## ☐ **Acknowledgment**

- ☐ UIUC CAD group
- ☐ EinsTimer team (Debjit, Kerim, Natesan, Hemlata, Adil, Jin, Michel, etc.)